

# A Description of the RCImmix Algorithm

Reference Counting with better heap allocation

Nathan Jervis

Department of Computer Science  
McMaster University, Hamilton  
jervisnd@mcmaster.ca  
1211159

March 26, 2014

# Overview

- Introduction to automatic memory management
- Problems with existing reference counting
- Optimizations to reference counting
- RCImmux algorithm

# Manual Memory Management

## Manual Memory Management

- Difficult to use
- Can cause dangling pointers
- Leads to memory leaks

*Much better if the compiler/runtime can manage memory for us*

# Automatic Memory Management

## Tracing Garbage Collector:

- Periodically pause program and follow program references
- Collect anything not referred to

## Reference Counting:

- Counter keeps track of how many things are pointing to it
- When counter reaches 0, free memory
- Allocate, Retain, Release

# Reference Counting vs Tracing

## **Tracing Garbage Collector:**

- Little work for allocation
- Better cache performance (with compacting)
- Requires pausing to collect

## **Reference Counting:**

- Doesn't pause
- Huge overhead
- Poor cache locality

# RCImmux Optimizations

## Optimizations

- Tracing collector as a backup
- Limited Bit Count
- Block Based allocation
- Nursery Allocation and Copying Collectors
- Objects are born as dead

# Tracing Backup Collector

- Naive reference counting won't collect cycles
- Reference counting handle normal memory collection
- When cycle garbage accumulates too much, call Tracer
- Tracer will very rarely be called, so you still don't have to worry about pauses too much

## Limited bit count

- In theory everything could point to object, requires 32/64 bits
- In practice most objects only a few things pointing to it
- Using 3-4 bits is fine
- If it overflows, just leave it at max, don't decrement
- Tracer will fix it when it collects



# Block Based Allocation

- Items are allocated out of a block
- Block keeps a pointer to the next free spot
- Count of number of live objects on the block
- Great for cache performance

# Nurseries and Copying Collection

- Objects are allocated into a nursery
- Move when mature (copied or pass collection cycle)
- Only new objects in nursery
  - Will be collected during collection cycles
  - Most objects die young, not copied
  - Nurseries are cheap way to collect

# Objects are born as dead

- When you create the object, garbage collector already considers it dead
- Only when it moves or matures do you consider it alive
- ModBuffer contains all objects that have created new objects
- Can process ModBuffer to check if "dead" objects are actually alive

## Typical Use

### Algorithm 1.1: USER CODE(*args*)

**main**

$x \leftarrow \text{ALLOCATE}(\textit{size})$

...

$\text{RETAIN}(x)$

$y \leftarrow x$

...

$\text{RELEASE}(y)$

...

$\text{RELEASE}(x)$

# Allocate

**Algorithm 2.2:** ALLOCATE(*size*)

**global** *bumpPointer*, *block*

**if** *size* < *block.end* − *bumpPointer*

**then**  $\begin{cases} \textit{pointer} \leftarrow \textit{bumpPointer} \\ \textit{bumpPointer} \leftarrow \textit{bumpPointer} + \textit{size} + 1 \end{cases}$

**else**  $\begin{cases} \textit{block} \leftarrow \text{GETNEWFREEBLOCK}() \\ \textbf{return} (\text{ALLOCATE}(\textit{size})) \end{cases}$

*pointer.new* ← **true**

*pointer.count* ← 0

**return** (*pointer*)

## Algorithm 2.3: RETAIN(*object*)

```
if object.new = false  
  then { if object.count not max  
        then object.count  $\leftarrow$  object.count + 1  
        return  
  pointer  $\leftarrow$  COPYTONEWLOCATION(object)  
  ADDTOMODBUFFER(object)  
  object.count  $\leftarrow$  2  
  object.block.liveCount  $\leftarrow$  object.block.liveCount + 1  
  return
```

## Algorithm 2.4: RELEASE(*object*)

```
if object.new = true
  then return
if object.count not max
  then object.count  $\leftarrow$  object.count - 1
if object.count = 0
  then { PROCESSMODBUFFER()
        object.block.liveCount  $\leftarrow$  object.block.liveCount - 1
        FREEBLOCKS() }
```

**Algorithm 3.5:** ADDTOMODBUFFER(*object*)

```
global ModBuffer  
MODBUFFER.PUSH(object)
```

**Algorithm 3.6:** PROCESSMODBUFFER(*none*)

```
global ModBuffer  
for each obj ∈ ModBuffer  
do { for each child ∈ obj.references  
    do { if child.new = true  
        then { RETAIN(child)child.count ← 1
```



## Conclusion

	Reference Counting	Tracing	RCImmix (Both)
Overhead	Significant	None	Minimal
Speed	Slow	Fast	Fast
Cache Performance	Poor	Good	Good
Pauses	Short	Long	Mostly short
Implementation Difficulty	Easy	Hard	Very Hard

## Resources



Xi Yang Kathryn S. McKinley Rifat Shahriyar, Stephen M. Blackburn.

Taking off the gloves with reference counting immix.  
pages 1–18, 2013.



Daniel Frampton Rifat Shahriyar, Stephen M. Blackburn.

Down for the count? getting reference counting back in the ring.  
pages 1–11, 2012.